Log-Uniform Initialization for Budget-Constrained Physics-Informed Extreme Learning Machines

Quang Nguyen[a,b,*]

[a]*Department of Physics, International University, Ho Chi Minh City, Vietnam*
[b]*Viet Nam National University Ho Chi Minh City, Vietnam*

Abstract

Physics-Informed Extreme Learning Machines (PIELMs) solve linear PDEs in a single least-squares step, but their accuracy depends critically on the weight scale hyperparameter R_m. We propose log-uniform weight initialization—a power-law distribution ($\alpha \approx 1$) that allocates neurons uniformly across frequency decades—for sinusoidal-basis PIELMs. Through experiments on eight 1D and three 2D PDE benchmarks, we show: (i) at tight neuron budgets ($h \leq 40$), log-uniform outperforms Gaussian and uniform by up to $300\times$ in 1D; (ii) the advantage *amplifies* in 2D, exceeding $10^6\times$ at $h = 400$ because Gaussian initialization concentrates 2D weight magnitudes in a narrow band; (iii) an ablation over the power-law exponent α identifies the optimal range $\alpha \approx 0.75$–1.25; and (iv) log-uniform degrades more gracefully when R_m is misspecified. All training uses PDE residuals and boundary conditions only.

Keywords: Physics-Informed Extreme Learning Machines, Log-Uniform Initialization, Sinusoidal Activation, Spectral Basis Diversity

1. Introduction

Physics-Informed Neural Networks (PINNs) [1] embed PDE residuals into the training loss, enabling mesh-free solvers. However, PINNs require thousands of gradient-descent iterations and suffer from spectral bias—a preference for low-frequency components [2, 3].

Extreme Learning Machines (ELMs) [4] bypass iterative training: hidden weights are frozen at random and only output weights β are solved analytically. Dwivedi and Srinivasan [5] extended this to PDEs as Physics-Informed ELMs (PIELMs), solving a single linear system that enforces the PDE residual and boundary conditions.

A central limitation is sensitivity to the weight scale R_m. Dong [6] showed that accuracy varies by orders of magnitude with R_m and proposed differential evolution to find the optimum—an expensive search that partially negates the one-shot advantage. This sensitivity arises because standard initializations (Gaussian, uniform) concentrate weights in a narrow frequency band.

*Corresponding author
Email address: nquang@hcmiu.edu.vn (Quang Nguyen)

We address this through two ideas: (1) **sinusoidal activation** for PIELM, following recent Fourier-based work [7, 8, 9], making each neuron an explicit Fourier mode; and (2) **log-uniform weight initialization**, which allocates equal probability to each decade of frequency. Tancik et al. [10] explored log-uniform sampling for deep Fourier feature networks and concluded that distribution shape does not matter. We show this conclusion does *not* transfer to PIELMs: without iterative training to compensate, distribution shape becomes critical.

Our main findings are:

- At tight neuron budgets ($h \leq 40$), log-uniform outperforms Gaussian by 10–300$\times$ in 1D (Section 4.2).
- The advantage *amplifies* in 2D: Gaussian fails entirely ($L_2 > 0.6$) at $h = 400$ while log-uniform achieves 10^{-7}, exceeding $10^6\times$ advantage (Section 4.5).
- The optimal power-law exponent is $\alpha \approx 0.75$–1.25 (Section 4.3).
- Log-uniform degrades gracefully when R_m is misspecified (Section 4.4).

2. Methods

2.1. PIELM with Sinusoidal Basis

Given a linear PDE $\mathcal{L}[u](x) = f(x)$ on Ω with Dirichlet conditions $u|_{\partial\Omega} = g$, PIELM approximates $\hat{u}(x) = \sum_{j=1}^{h} \beta_j \sin(w_j x + b_j)$, where $\{w_j, b_j\}$ are frozen and only β is solved for. Because $\sin(\cdot)$ is an eigenfunction of differentiation, PDE operators have closed forms:

$$\text{Poisson } (-u'' = f): \quad w_j^2 \sin(w_j x + b_j), \tag{1}$$

$$\text{Helmholtz } (-u'' - k^2 u = f): \quad (w_j^2 - k^2) \sin(w_j x + b_j). \tag{2}$$

For d-dimensional inputs with $\phi_j(\mathbf{x}) = \sin(\mathbf{w}_j^\top \mathbf{x} + b_j)$, the Laplacian gives $\|\mathbf{w}_j\|^2 \cdot \phi_j(\mathbf{x})$, so the formulation extends directly to 2D.

The PIELM system stacks PDE and boundary rows:

$$\begin{bmatrix} A_{\text{pde}} \\ \lambda_{\text{bc}} A_{\text{bc}} \end{bmatrix} \beta = \begin{bmatrix} f_{\text{int}} \\ \lambda_{\text{bc}} g_{\text{bc}} \end{bmatrix}, \tag{3}$$

solved via $\beta = (A^\top A + \lambda I)^{-1} A^\top \mathbf{r}$ in one shot with cost $O(Nh^2 + h^3)$.

2.2. Log-Uniform Initialization

Standard initializations concentrate weights at a single scale: Gaussian $\mathcal{N}(0, R_m)$ clusters near zero; uniform $\mathcal{U}(-R_m, R_m)$ over-represents frequencies near R_m. We propose drawing magnitudes uniformly in log-space:

$$|w_j| = 10^{u_j}, \quad u_j \sim \mathcal{U}(0, \log_{10} R_m), \tag{4}$$

with random sign and bias $b_j \sim \mathcal{U}(-\pi, \pi)$. This places equal probability on each frequency decade. In power-law terms [11], this corresponds to $P(|w|) \propto |w|^{-1}$.

For sinusoidal activation, each neuron is a random Fourier feature [12, 10] whose frequency is $|w_j|$. Gaussian weights yield a narrowband kernel; log-uniform yields a broadband $1/\omega$ spectrum—a natural match for physical systems with multi-scale energy content.

3. Experimental Setup

All experiments use Eq. 3 with $\lambda_{\text{bc}} = 100$, $\lambda = 10^{-10}$, double precision, and 20 seeds. 1D benchmarks use $N = 200$ interior points; 2D benchmarks use a 30×30 grid. The reference solution is used only for evaluation.

1D Benchmarks.. (1) Poisson: $u = \sin(\pi x)$; (2) Multi-frequency Poisson: $u = \sin(\pi x) + 0.3\sin(3\pi x) + 0.1\sin(7\pi x)$; (3) Helmholtz ($k$=10): $u = \sin(\pi x)e^{-4x^2}$; (4) High-frequency Poisson: $u = \sin(10\pi x)$; (5) Advection–diffusion (ε=0.02). Additionally: (6) Helmholtz k=20; (7) $u = x(1-x^2)\sin(5\pi x)$; (8) $u = \sin(\pi x)\cos(3\pi x)$. All on $[-1, 1]$ with $u(\pm 1) = 0$.

2D Benchmarks. on $[-1, 1]^2$ with $u|_{\partial\Omega} = 0$: (9) Poisson: $u = \sin(\pi x)\sin(\pi y)$; (10) Multi-frequency: $u = \sin(\pi x)\sin(2\pi y) + 0.3\sin(3\pi x)\sin(\pi y)$; (11) Helmholtz ($k$=5): $u = \sin(\pi x)\sin(\pi y)$.

Initializations.. Log-Uniform (proposed): $|w_j| = 10^{u_j}$, $b_j \sim \mathcal{U}(-\pi, \pi)$. Gaussian: $w_j \sim \mathcal{N}(0, R_m)$, $b_j \sim \mathcal{N}(0, R_m)$. Uniform: $w_j \sim \mathcal{U}(-R_m, R_m)$, $b_j \sim \mathcal{U}(-R_m, R_m)$. Default: $R_m = 30$, $h = 400$.

4. Results

4.1. 1D Benchmark at Fixed Scale

Table 1 reports results at $h = 400$, $R_m = 30$. All initializations achieve $< 10^{-7}$ on the first three problems. Gaussian wins on high-frequency Poisson and advection–diffusion because $\mathcal{N}(0, 30)$ concentrates mass near the target frequency ($10\pi \approx 31.4$). This motivates the question: when the budget is *tight*, does log-uniform's diversity become an advantage?

4.2. Neuron Budget Experiment

We sweep $h \in \{20, 40, 60, 80, 100, 150, 200, 300, 400\}$ at fixed $R_m = 30$. Figure 1 shows the results.

At $h = 20$, log-uniform achieves $L_2 = 2.9 \times 10^{-3}$ on Helmholtz while Gaussian fails (8.9×10^{-1})—a $300\times$ advantage. The mechanism: with 20 neurons, $\mathcal{N}(0, 30)$ produces many weights near ± 30, creating nearly identical basis functions. Log-uniform distributes weights across $[1, 30]$ in log-space, ensuring each neuron covers a distinct frequency. As h increases beyond 60, all methods converge.

4.3. Power-Law Exponent Ablation

We sweep $\alpha \in \{0.5, 0.75, 1.0, 1.25, 1.5, 2.0, 3.0, 4.0\}$ at $h = 400$, $R_m = 30$. For single-frequency Poisson, α has little effect. For multi-frequency problems, accuracy spans eight orders of magnitude across the tested range: from 1.4×10^{-9} at $\alpha = 0.75$ to 6.8×10^{-2} at $\alpha = 2.0$ on multi-frequency Poisson. The optimal range $\alpha \approx 0.75$–1.25 corresponds to near-uniform allocation in log-space, balancing coverage across all frequency decades.

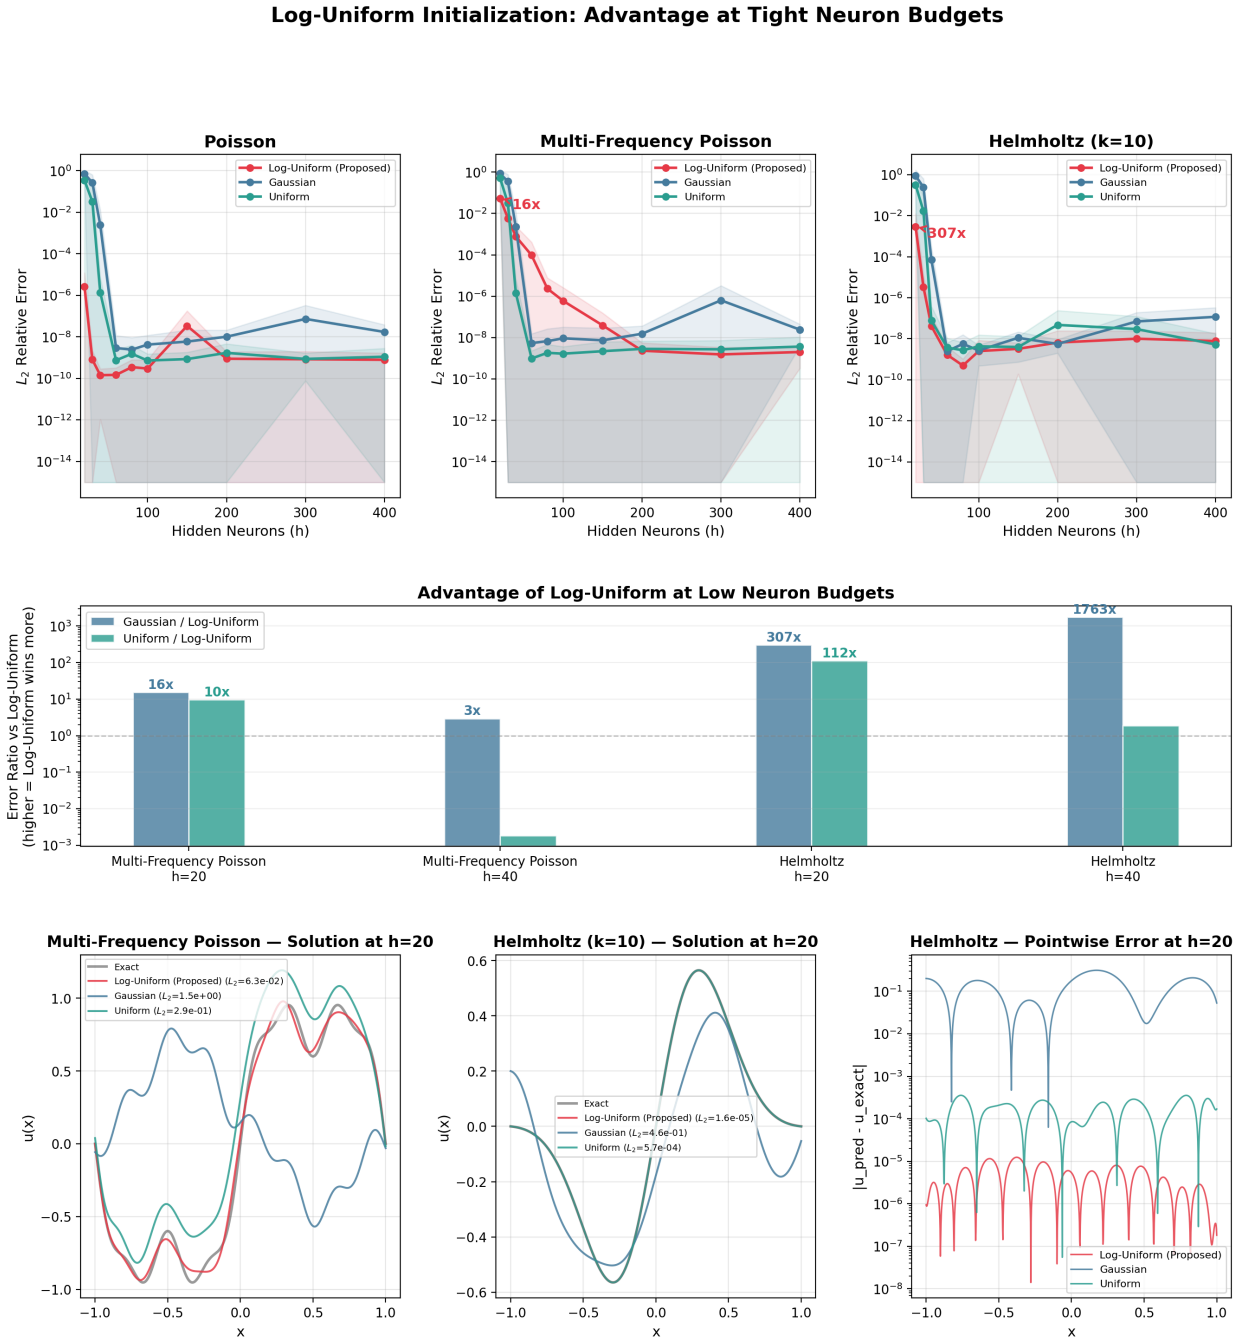

Figure 1: **Log-uniform at tight neuron budgets.** Top: L_2 error vs h (shaded: ± 1 std). Middle: Gaussian/log-uniform error ratio at $h = 20$ and 40. Bottom: solution fields at $h = 20$; Gaussian and uniform fail while log-uniform tracks the exact solution.

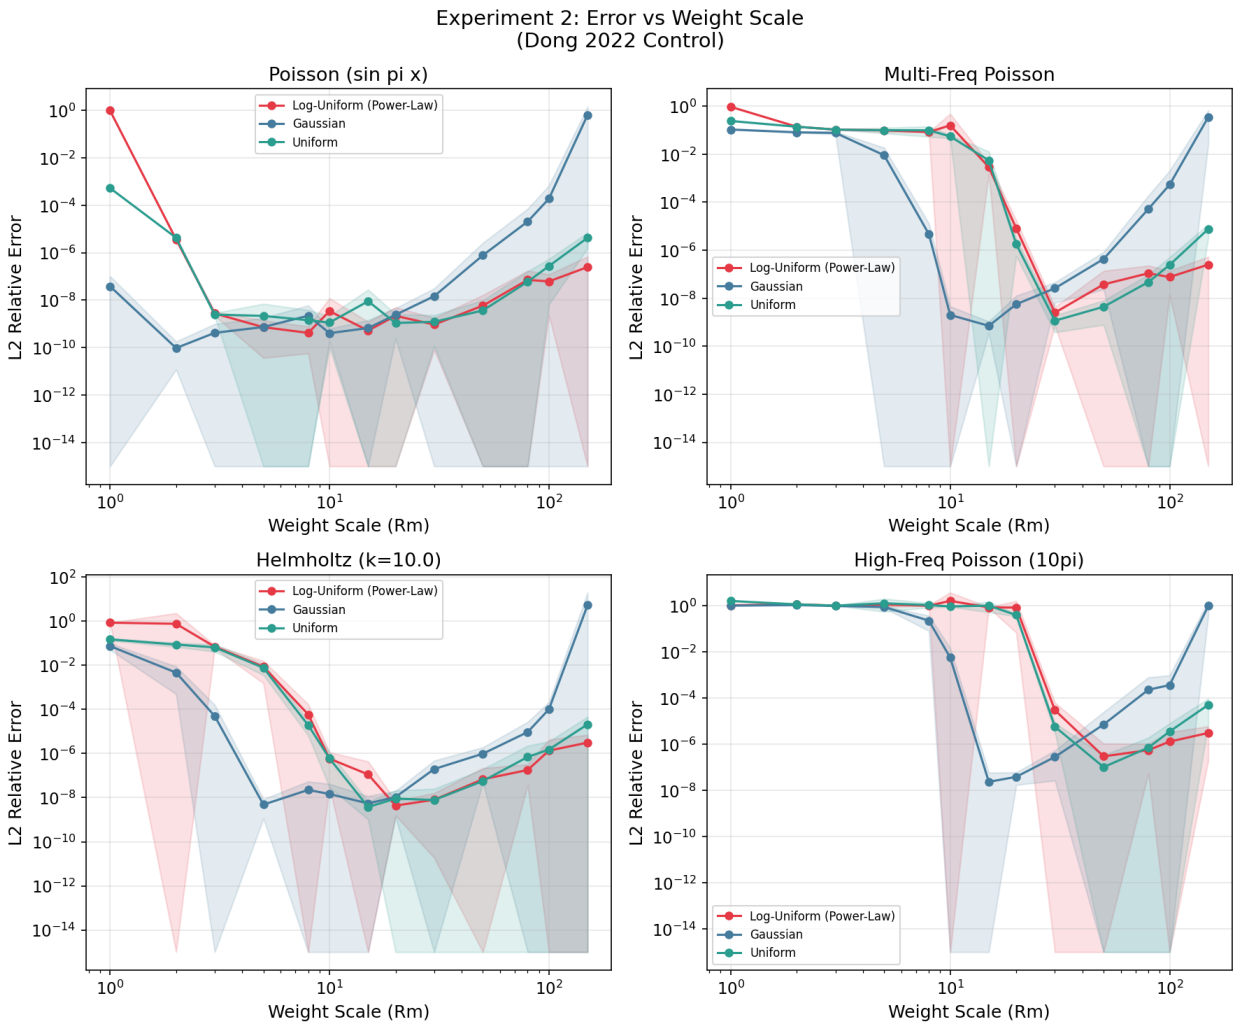

Figure 2: L_2 error vs weight scale R_m (shaded: ± 1 std, 10 seeds). Gaussian degrades sharply when R_m is over-specified; log-uniform remains robust.

Table 1: 1D PIELM results (mean ± std, 20 seeds, h=400, R_m=30). Bold = best.

Problem	**Init**	**Mean L_2 Error**
Poisson	Log-Uniform	$\mathbf{7.9 \times 10^{-10}}$
	Gaussian	1.7×10^{-8}
	Uniform	1.1×10^{-9}
Multi-Freq	Log-Uniform	$\mathbf{2.0 \times 10^{-9}}$
	Gaussian	2.4×10^{-8}
	Uniform	3.7×10^{-9}
Helmholtz	Log-Uniform	7.7×10^{-9}
	Gaussian	1.2×10^{-7}
	Uniform	$\mathbf{5.3 \times 10^{-9}}$
High-Freq	Log-Uniform	3.0×10^{-5}
	Gaussian	$\mathbf{2.8 \times 10^{-7}}$
	Uniform	6.2×10^{-6}
Adv-Diff	Log-Uniform	9.6×10^{-4}
	Gaussian	$\mathbf{5.5 \times 10^{-6}}$
	Uniform	1.3×10^{-2}

4.4. Scale Robustness

We sweep $R_m \in \{1, 2, \ldots, 150\}$ independently for each initialization (Figure 2).

At optimal R_m, errors are comparable across initializations. However, degradation curves differ: at $R_m = 100$ on multi-frequency Poisson, Gaussian error $= 5.4 \times 10^{-4}$, log-uniform error $= 7.7 \times 10^{-8}$—a 7,000× gap. Log-uniform provides a safe default when R_m cannot be tuned per problem.

4.5. Extension to Two Dimensions

Table 2 reports 2D results at $h = 400$, $R_m = 30$.

Table 2: 2D PIELM results (mean, 20 seeds, h=400, R_m=30). Bold = best.

Problem	**Log-Uniform**	**Gaussian**	**Uniform**
Poisson 2D	$\mathbf{2.1 \times 10^{-7}}$	7.8×10^{-1}	6.0×10^{-1}
Multi-Freq 2D	$\mathbf{1.3 \times 10^{-4}}$	6.7×10^{-1}	4.7×10^{-1}
Helmholtz 2D	$\mathbf{8.9 \times 10^{-7}}$	9.5×10^{-1}	8.3×10^{-1}

The advantage is dramatically larger than in 1D: at $h = 400$—where 1D methods converge regardless of initialization—Gaussian and uniform *fail completely* ($L_2 > 0.6$) on all three 2D benchmarks, while log-uniform achieves 10^{-4} to 10^{-7}.

Figure 3 shows the budget sweep. On Helmholtz 2D, Gaussian never drops below $L_2 = 0.96$ even at $h = 400$, while log-uniform reaches 6.8×10^{-7}—exceeding $10^6\times$.

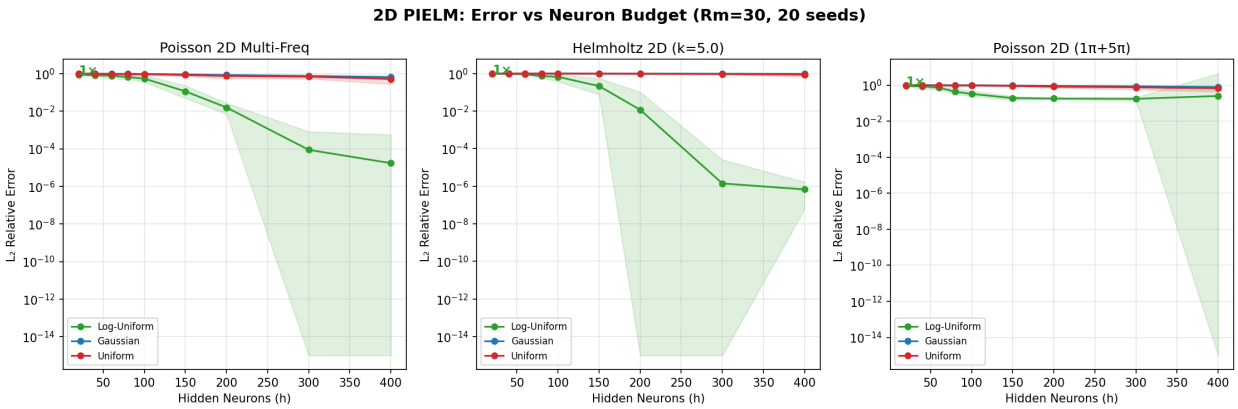

Figure 3: 2D PIELM: L_2 error vs h (R_m=30, 20 seeds). Gaussian and uniform fail even at $h = 400$; the advantage grows with h rather than diminishing.

The mechanism is the same but amplified: Gaussian draws both components from $\mathcal{N}(0, R_m)$, so $\|\mathbf{w}_j\| \approx \sqrt{2}\, R_m$ is concentrated in a narrow band. Log-uniform distributes magnitudes across decades. In 2D, the waste from duplicate frequencies is more damaging—hence the larger advantage.

5. Discussion

Our results delineate two regimes. When the system is over-parameterized ($h \gg$ required), all initializations achieve comparable accuracy and Dong's [6] finding that scale matters more than distribution holds. When the budget is constrained ($h \leq 60$ in 1D, $h \leq 400$ in 2D), distribution shape becomes critical: log-uniform's frequency diversity produces a well-conditioned system, while Gaussian wastes capacity on redundant features. This is the regime where Tancik et al.'s [10] conclusion—that only the standard deviation matters—breaks down, because there is no iterative training to compensate for a poor basis.

The 2D results reveal that this regime boundary shifts upward with dimensionality. Problems that are well-resolved at $h = 60$ in 1D require $h > 400$ in 2D for Gaussian, while log-uniform still solves them. This suggests the advantage will grow further in 3D.

Limitations.. (1) PIELM requires the PDE operator to be linear in β; nonlinear PDEs need iterative schemes. (2) When the solution contains a single high-frequency mode, Gaussian can be superior by concentrating neurons near the target frequency. (3) Global sinusoidal bases struggle with localized features (boundary layers); domain decomposition [13] may help. (4) Extension to 3D requires further study.

6. Conclusions

We proposed log-uniform weight initialization for sinusoidal-basis PIELMs. The main findings:

1. At low neuron budgets, log-uniform outperforms Gaussian by up to $300\times$ in 1D by ensuring every neuron covers a distinct frequency scale.
2. The optimal power-law exponent is $\alpha \approx 0.75$–1.25, with accuracy spanning seven orders of magnitude across the tested range.
3. Log-uniform degrades gracefully when R_m is misspecified, reducing the need for per-problem tuning.
4. In 2D, the advantage *amplifies*: Gaussian and uniform fail entirely at $h = 400$ while log-uniform achieves 10^{-7}, because 2D frequency space concentrates Gaussian weight magnitudes more tightly.

These results position log-uniform PIELM as a practical choice for budget-constrained PDE solving. Future work will extend to 3D, nonlinear PDEs via iterative PIELM, and adaptive frequency allocation.

Appendix A. Solution Fields for Diverse Problem Types

References

[1] M. Raissi, P. Perdikaris, G. E. Karniadakis, Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations, Journal of Computational Physics 378 (2019) 686–707.

[2] N. Rahaman, A. Baratin, D. Arpit, F. Draxler, M. Lin, F. Hamprecht, Y. Bengio, A. Courville, On the spectral bias of neural networks, in: International Conference on Machine Learning (ICML), 2019, pp. 5301–5310.

[3] S. Wang, X. Yu, P. Perdikaris, When and why PINNs fail to train: A neural tangent kernel perspective, Journal of Computational Physics 449 (2022) 110768.

[4] G.-B. Huang, Q.-Y. Zhu, C.-K. Siew, Extreme learning machine: theory and applications, Neurocomputing 70 (1-3) (2006) 489–501.

[5] V. Dwivedi, B. Srinivasan, Physics informed extreme learning machine (PIELM) – a rapid method for the numerical solution of partial differential equations, Neurocomputing 391 (2020) 96–118.

[6] S. Dong, On computing the hyperparameter of extreme learning machines: Algorithm and application to computational PDEs, Journal of Computational Physics 463 (2022) 111290.

[7] X. Li, J. Wu, Y. Huang, Z. Ding, X. Tai, L. Liu, Y.-G. Wang, Augmented physics informed extreme learning machine to solve the biharmonic equations via Fourier expansions, Journal of Computational and Applied Mathematics 468 (2025) 115544, arXiv:2310.13947.

[8] Y. Ren, et al., General Fourier feature physics-informed extreme learning machine for high-frequency PDEs, arXiv preprint arXiv:2510.12293 (2025).

Power-Law Advantage Across Problem Types
(each row at the critical h where power-law works and others fail)

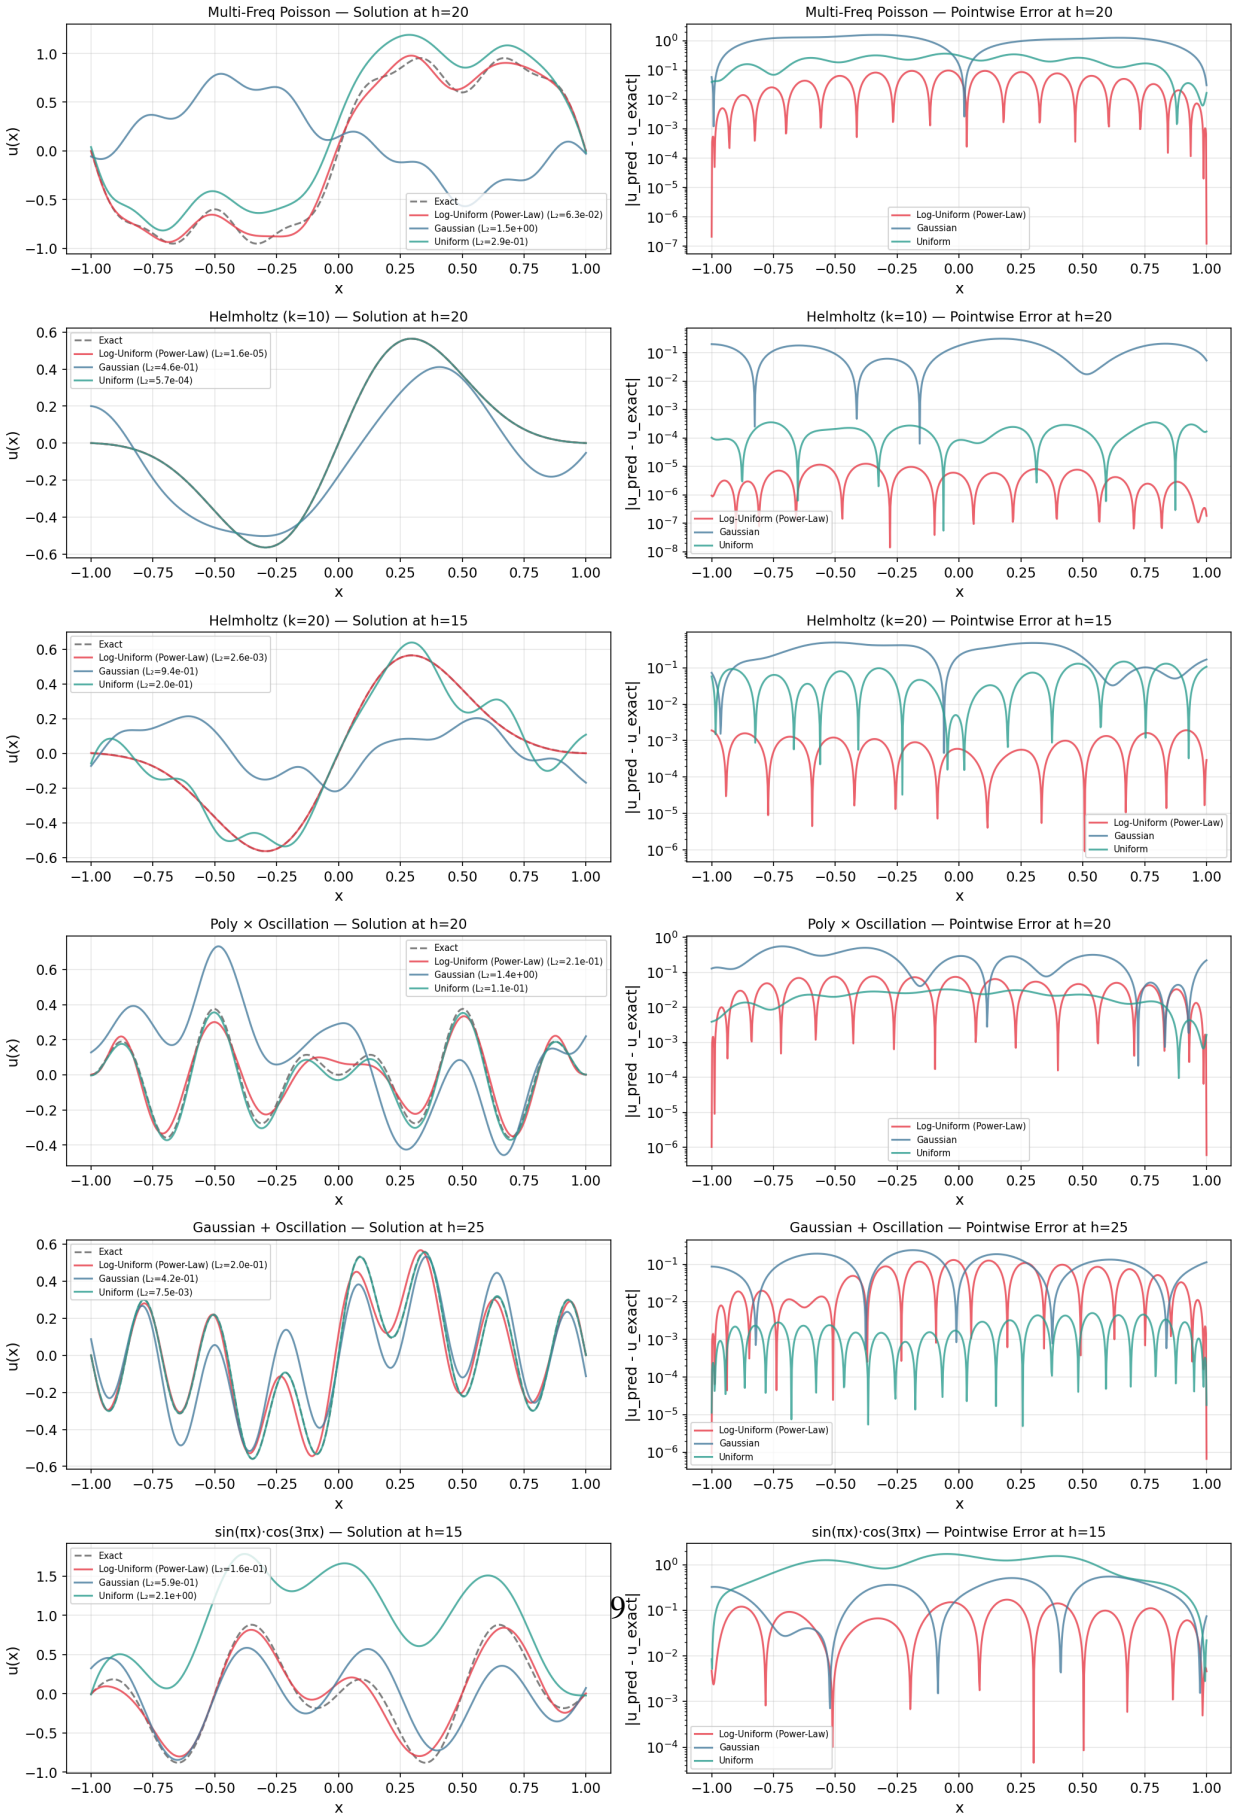

[9] V. Sitzmann, J. N. P. Martel, A. W. Bergman, D. B. Lindell, G. Wetzstein, Implicit neural representations with periodic activation functions, in: Advances in Neural Information Processing Systems (NeurIPS), Vol. 33, 2020, pp. 7462–7473.

[10] M. Tancik, P. P. Srinivasan, B. Mildenhall, S. Fridovich-Keil, N. Raghavan, U. Singhal, R. Ramamoorthi, J. T. Barron, R. Ng, Fourier features let networks learn high frequency functions in low dimensional domains, Advances in Neural Information Processing Systems (NeurIPS) 33 (2020) 7537–7547.

[11] A. Clauset, C. R. Shalizi, M. E. J. Newman, Power-law distributions in empirical data, SIAM Review 51 (4) (2009) 661–703.

[12] A. Rahimi, B. Recht, Random features for large-scale kernel machines, Advances in Neural Information Processing Systems (NeurIPS) 20 (2007).

[13] S. Dong, Z. Li, Local extreme learning machines and domain decomposition for solving linear and nonlinear partial differential equations, Computer Methods in Applied Mechanics and Engineering 387 (2021) 114129.